

## Numerical Analysis (CS 450)

# Homework Set 1, Bill Karr

### Problem 1: (15 points)

Consider the function  $f : \mathbb{R}^2 \rightarrow \mathbb{R}$  defined by  $f(x, y) = x - y$ . Measuring the size of the input  $(x, y)$  by  $|x| + |y|$ , and assuming that  $|x| + |y| \approx 1$  and  $x - y \approx \varepsilon$ , show that  $\text{cond}(f) \approx \frac{1}{\varepsilon}$ . What can you conclude about the sensitivity of subtraction?

*Solution & Proof.* To compute the condition number of  $f$ , we need to examine

$$\text{cond}(f) = \frac{\frac{|f(x + \Delta x, y + \Delta y) - f(x, y)|}{|(\Delta x, \Delta y)|}}{\frac{|f(x, y)|}{|(x, y)|}} = \frac{\frac{|\Delta x - \Delta y|}{|\Delta x| + |\Delta y|}}{\frac{|x - y|}{|x| + |y|}} \approx \frac{\frac{|\Delta x - \Delta y|}{\varepsilon}}{\frac{1}{1}} = \frac{1}{\varepsilon} \frac{|\Delta x - \Delta y|}{|\Delta x| + |\Delta y|}$$

Using the triangle inequality, we see that  $|\Delta x - \Delta y| \leq |\Delta x| + |\Delta y|$  and equality can be achieved, for example, if  $\Delta x > 0$  and  $\Delta y < 0$ . Thus,  $\text{cond}(f) \approx \frac{1}{\varepsilon}$  for  $x - y \approx \varepsilon$ .

Thus, subtraction of two numbers is very sensitive when the two numbers are relatively large, but have relatively small difference. It's easy to have a large relative error.  $\square$

### Problem 2: (15 points)

For computing the midpoint  $m$  of an interval  $[a, b]$ , which of the following two formulae is preferable in floating-point arithmetic?

(i)  $m = (a + b)/2$

(ii)  $m = a + (b - a)/2$

Why? When? Devise examples for which the “midpoint” given by the formula lies outside of the interval  $[a, b]$ . Specify the floating point system used, specify  $a$  and  $b$  in floating-point representation. Write down all intermediate results for both cases. Point out the steps where the problem occurs.

*Solution.* Neither formula is the “preferable” formula. They both work in certain circumstances. Sometimes one works and the other doesn't. Let us work with in a binary floating-point system ( $\beta = 2$ ) with  $p = 2$ ,  $U = 4$ , and  $L = -4$  where anything that overflows is considered infinite.

Let  $a = -(0.1)_2 \times 2^4$  and  $b = (1.1)_2 \times 2^4$ . Then,  $a + b = (1.0)_2 \times 2^4$ . We divide by two and obtain  $m = (1.0)_2 \times 2^3$ . However,  $b - a = (1.1)_2 \times 2^5$  which is larger than overflow and thus infinite. We compute  $m = a + (b - a)/2 = a + \infty = \infty$  which is larger than  $b$ . Thus, the second formula fails.

On the other hand, if  $a = (0.1)_2 \times 2^4$  and  $b = (1.1)_2 \times 2^4$ , then  $a + b = (1.1)_2 \times 2^5$  which is infinite and  $m = \infty$  using the first formula which is larger than  $b$ . However, the second formula doesn't have the same issue. We get  $m = a + (b - a)/2 = (0.1)_2 \times 2^4 + (1.0)_2 \times 2^4/2 = (1.0)_2 \times 2^4$ .  $\square$

### Problem 3: Bessel recurrence vs. floating point (20 points)

- (a) Write a program that tests the accuracy to which the values returned from `scipy` obey

$$J_{n+1}(z) = (2n/z)J_n(z) - J_{n-1}(z). \quad (1)$$

This should yield roughly machine precision in each case.

- (b) Write a program that uses the values for  $J_0(z)$  and  $J_1(z)$  (obtained from `scipy`) and the so-called ‘recurrence relation’ (1) to compute the values of  $J_2(z), \dots, J_{50}(z)$ .

Print these values for each  $n$ , and also print the relative error compared to the value returned by `scipy` and report your results.

For your experiments, fix  $z = 20$ .

- (c) You should find that the results from (1) rapidly start losing precision around  $n = 30$ . Identify the reason for this loss of precision.
- (d) Observe that (1) can be rearranged to compute  $J_{n-1}$  from  $J_{n+1}$ :

$$J_{n-1}(z) = (2n/z)J_n(z) - J_{n+1}(z). \quad (2)$$

Do you believe using (2) to compute  $J_0, \dots, J_{48}$  from  $J_{49}$  and  $J_{50}$  will encounter loss of precision? Why?

*Solution.* (a) See Table 1.

- (b) See Table 2.

- (c) On Table 2, we see that the accuracy blows up for the larger values of  $n$ . As  $n$  grows, we see that  $J_n(20)$  becomes smaller in magnitude. thus, the recurrence relation is using two larger numbers to compute a smaller number by taking a difference. We showed in the first problem that when  $|x| + |y|$  is much bigger than  $|x - y|$ , subtraction becomes sensitive and prone to error. All of this error continues to propagate as we use the recurrence relation to compute  $J_n(z)$  for larger values of  $n$ .

- (d) See Table 3. Here, since we started with small numbers for  $J_{50}(20)$  and  $J_{49}(20)$ , the actual error starts out small relative to the size of  $J_n(20)$  for smaller  $n$  values, thus the propagated error is small relative to the values of the Bessel function that we’re computing.

□

Table 1: Bessel Recursion Accuracy

$n$	$J_n(20)$	$J_n(z) - ((2n/z)J_{n-1}(z) - J_{n-2}(z))$	% error
2	0.167024664341	2.77555756156e-17	1.73103040998e-14
3	0.0668331241758	0.0	0.0
4	-0.160341351923	0.0	0.0
5	-0.0989013945604	0.0	0.0
6	0.130670933555	-6.93889390391e-18	1.25964630952e-14
7	0.151169767982	0.0	0.0
8	-0.0550860495637	1.38777878078e-17	1.87870435183e-14
9	-0.184221397721	-2.77555756156e-17	2.21820557913e-14
10	-0.0738689288408	0.0	0.0
11	0.125126254648	6.93889390391e-18	1.13091785556e-14
12	0.186482558024	-2.77555756156e-17	2.33258509034e-14
13	0.061356303376	-2.77555756156e-17	1.35960069907e-14
14	-0.11899062431	0.0	0.0
15	-0.204145052548	-1.8323016715e-17	2.25633726574e-12
16	-0.146397944003	-2.77555756156e-17	1.91180645573e-14
17	-0.000812069055154	2.77555756156e-17	1.19071633614e-14
18	0.14517984042	5.55111512313e-17	2.21080831413e-14
19	0.233099813727	5.55111512313e-17	2.53635513253e-14
20	0.251089842916	0.0	0.0
21	0.218861903522	2.77555756156e-17	2.50878255518e-14
22	0.164747773775	-2.77555756156e-17	4.10689455014e-14
23	0.110633644029	1.38777878078e-17	3.64737607094e-14
24	0.0675828786855	-1.38777878078e-17	6.96357762909e-14
25	0.0380486890792	1.73472347598e-18	1.77353447714e-14
26	0.0199291061966	1.73472347598e-18	3.83465294446e-14
27	0.00978116579257	-8.67361737988e-19	4.37898765054e-14
28	0.00452380828487	-3.25260651746e-19	3.94648436394e-14
29	0.00198073574809	1.62630325873e-19	4.97396255084e-14
30	0.000824178234983	1.89735380185e-19	1.52993447482e-13
31	0.000326963309857	4.06575814682e-20	9.01842801363e-14
32	0.000124015363604	3.38813178902e-20	2.15238955394e-13
33	4.50827809534e-05	4.23516473627e-21	8.00712895667e-14
34	1.57412573519e-05	-1.05879118407e-21	6.18004041909e-14
35	5.2892425727e-06	1.05879118407e-21	1.97615269075e-13
36	1.71324313802e-06	-5.29395592034e-23	3.26784982006e-14
37	5.35784096556e-07	1.19114008208e-22	2.51188206646e-13
38	1.62001199928e-07	2.31610571515e-23	1.72154758836e-13
39	4.74202231856e-08	8.27180612553e-25	2.23347728163e-14
40	1.34536258586e-08	2.06795153138e-24	2.08833590054e-13
41	3.7035550769e-09	1.55096364854e-25	6.02548368075e-14
42	9.90238941374e-10	7.75481824268e-26	1.19114529287e-13
43	2.57400688594e-10	3.23117426779e-27	2.01499862811e-14
44	6.51038818615e-11	-5.65455496862e-27	1.46899655411e-13
45	1.60356152243e-11	-7.06819371078e-28	7.84357303505e-14
46	3.84926360297e-12	4.79627430374e-28	2.32956380393e-13
47	9.01144628754e-13	3.78653234506e-29	8.24295069619e-14
48	2.05887226426e-13	1.4199496294e-29	1.41784359854e-13
49	4.59366128055e-14	3.54987407349e-30	1.66294964838e-13
50	1.001485376e-14	1.38050658414e-30	3.10153763163e-13

Table 2: Estimating  $J_{n+1}(z)$  using  $J_0(z)$  and  $J_1(z)$ 

$n$	$J_n(20)$ (scipy)	(using recurrence relation)	Relative error
0	0.167024664341	0.167024664341	0.0
1	0.0668331241758	0.0668331241758	0.0
2	-0.160341351923	-0.160341351923	1.73103040998e-16
3	-0.0989013945604	-0.0989013945604	1.40319435024e-16
4	0.130670933555	0.130670933555	2.1240818337e-16
5	0.151169767982	0.151169767982	1.8360533317e-16
6	-0.0550860495637	-0.0550860495637	1.25964630952e-16
7	-0.184221397721	-0.184221397721	1.5066423314e-16
8	-0.0738689288408	-0.0738689288408	3.75740870366e-16
9	0.125126254648	0.125126254648	2.21820557913e-16
10	0.186482558024	0.186482558024	2.97674762828e-16
11	0.061356303376	0.061356303376	3.39275356668e-16
12	-0.11899062431	-0.11899062431	0.0
13	-0.204145052548	-0.204145052548	0.0
14	-0.146397944003	-0.146397944003	0.0
15	-0.000812069055154	-0.000812069055154	2.25633726574e-14
16	0.14517984042	0.14517984042	1.91180645573e-16
17	0.233099813727	0.233099813727	0.0
18	0.251089842916	0.251089842916	2.21080831413e-16
19	0.218861903522	0.218861903522	7.60906539759e-16
20	0.164747773775	0.164747773775	1.68473145218e-15
21	0.110633644029	0.110633644029	3.76317383278e-15
22	0.0675828786855	0.0675828786855	8.62447855529e-15
23	0.0380486890792	0.0380486890792	2.2978469247e-14
24	0.0199291061966	0.0199291061966	7.10284918167e-14
25	0.00978116579257	0.00978116579257	2.5840397332e-13
26	0.00452380828487	0.00452380828487	1.0840563874e-12
27	0.00198073574809	0.00198073574808	5.16063694616e-12
28	0.000824178234983	0.00082417823496	2.75365946494e-11
29	0.000326963309857	0.000326963309804	1.63089434293e-10
30	0.000124015363604	0.000124015363472	1.06394506542e-09
31	4.50827809534e-05	4.50827806109e-05	7.59740964285e-09
32	1.57412573519e-05	1.57412564221e-05	5.90704255823e-08
33	5.2892425727e-06	5.28923993972e-06	4.97799916786e-07
34	1.71324313802e-06	1.71323537901e-06	4.52884120459e-06
35	5.35784096556e-07	5.35760348919e-07	4.43231451601e-05
36	1.62001199928e-07	1.61925842207e-07	0.00046516767109
37	4.74202231856e-08	4.71726830267e-08	0.00522013905183
38	1.34536258586e-08	1.26130849915e-08	0.0624769022028
39	3.7035550769e-09	7.57039941061e-10	0.795591013137
40	9.90238941374e-10	-9.66062922138e-09	10.7558567107
41	2.57400688594e-10	-3.93995568266e-08	154.067021855
42	6.51038818615e-11	-1.51877553768e-07	2333.84943117
43	1.60356152243e-11	-5.98486168997e-07	37323.3079143
44	3.84926360297e-12	-2.42161297292e-06	629111.713814
45	9.01144628754e-13	-1.00566109119e-05	11159819.958
46	2.05887226426e-13	-4.28331361304e-05	208041737.605
47	4.59366128055e-14	-0.000186975815288	4070300440.42
48	1.001485376e-14	-0.000835953195723	83471333258.1
49	2.13468524255e-15	-0.00382559952418	1.79211410091e+12
50	4.4510392847e-16	-0.0179094844728	4.02366353726e+13

Table 3: Estimating  $J_n(z)$  using  $J_{50}(z)$  and  $J_{49}(z)$ 

$n$	$J_n(20)$ (scipy)	(using recurrence relation)	Relative error
0	0.167024664341	0.167024664341	3.15735366857e-15
1	0.0668331241758	0.0668331241758	1.2458900863e-15
2	-0.160341351923	-0.160341351923	3.11585473797e-15
3	-0.0989013945604	-0.0989013945604	1.82415265531e-15
4	0.130670933555	0.130670933555	3.39853093392e-15
5	0.151169767982	0.151169767982	2.38686933121e-15
6	-0.0550860495637	-0.0550860495637	4.91262060714e-15
7	-0.184221397721	-0.184221397721	2.86262042967e-15
8	-0.0738689288408	-0.0738689288408	1.1272226111e-15
9	0.125126254648	0.125126254648	3.54912892661e-15
10	0.186482558024	0.186482558024	2.53023548403e-15
11	0.061356303376	0.061356303376	5.6545892778e-16
12	-0.11899062431	-0.11899062431	3.84876539906e-15
13	-0.204145052548	-0.204145052548	2.85516146804e-15
14	-0.146397944003	-0.146397944003	2.08548920446e-15
15	-0.000812069055154	-0.000812069055154	1.48330810784e-13
16	0.14517984042	0.14517984042	3.25007097474e-15
17	0.233099813727	0.233099813727	2.85771920674e-15
18	0.251089842916	0.251089842916	2.65296997696e-15
19	0.218861903522	0.218861903522	2.53635513253e-15
20	0.164747773775	0.164747773775	2.35862403306e-15
21	0.110633644029	0.110633644029	2.38334342742e-15
22	0.0675828786855	0.0675828786855	2.25879200258e-15
23	0.0380486890792	0.0380486890792	2.18842564257e-15
24	0.0199291061966	0.0199291061966	2.08907328873e-15
25	0.00978116579257	0.00978116579257	2.12824137257e-15
26	0.00452380828487	0.00452380828487	2.10905911945e-15
27	0.00198073574809	0.00198073574809	1.97054444274e-15
28	0.000824178234983	0.000824178234983	1.97324218197e-15
29	0.000326963309857	0.000326963309857	1.65798751695e-15
30	0.000124015363604	0.000124015363604	1.52993447482e-15
31	4.50827809534e-05	4.50827809534e-05	1.20245706848e-15
32	1.57412573519e-05	1.57412573519e-05	1.29143373237e-15
33	5.2892425727e-06	5.2892425727e-06	1.4412832122e-15
34	1.71324313802e-06	1.71324313802e-06	1.23600808382e-15
35	5.35784096556e-07	5.35784096556e-07	1.38330688352e-15
36	1.62001199928e-07	1.62001199928e-07	9.80354946017e-16
37	4.74202231856e-08	4.74202231856e-08	9.76843025847e-16
38	1.34536258586e-08	1.34536258586e-08	9.83741479063e-16
39	3.7035550769e-09	3.7035550769e-09	8.93390912652e-16
40	9.90238941374e-10	9.90238941374e-10	8.35334360215e-16
41	2.57400688594e-10	2.57400688594e-10	8.033978241e-16
42	6.51038818615e-11	6.51038818615e-11	5.95572646434e-16
43	1.60356152243e-11	1.60356152243e-11	6.04499588433e-16
44	3.84926360297e-12	3.84926360297e-12	6.29569951762e-16
45	9.01144628754e-13	9.01144628754e-13	4.48204173432e-16
46	2.05887226426e-13	2.05887226426e-13	3.67825863779e-16
47	4.59366128055e-14	4.59366128055e-14	2.74765023206e-16
48	1.001485376e-14	1.001485376e-14	1.57538177616e-16
49	2.13468524255e-15	2.13468524255e-15	0.0
50	4.4510392847e-16	4.4510392847e-16	0.0

Unfortunately, I had a difficult time getting used to Python because I have almost no programming experience and ran out of time to complete the rest of this assignment. I'm getting a study group together so that this doesn't happen again.